


(a) Model Comparison (clean test, $n = 40,230$)



(b) Adversarial Detection: Clean vs Robust ($n = 50,000$ / variant)

